


Low-precision GEMM: BF16 versus FP8 versus NVFP4

Three repeated hot-cache measurements per matched configuration.



FP32 accumulation/output; tile 256x128; cluster 2x1; NVFP4 uses one E4M3 scale per 16 values.